

2 Updated Table 3 as supporting material for
3 COLM rebuttal, new results are marked in
4 red.

Method	TNR	TPR	F1 ¹	TFLOPs
<i>Qwen3-8B-thinking</i>				
Basic Prompt	12.8	94.0	22.4	72.64
Probe-only	18.0	80.4	29.3	65.01
Prompt Opt.	18.0	78.6	29.1	69.59
+Cons@4	18.7	88.9	30.7	274.33
+Cons@8	18.8	92.1	31.2	545.19
VerifySteer-Uni	21.5	75.6	33.4	77.72
VerifySteer-Bi	21.5	76.2	33.4	77.71
<i>FARE-20B</i>				
Probe-only	26.6	90.5	41.1	5.85
FARE Prompt	28.5	76.2	41.5	7.73
+Cons@4	27.4	85.7	41.5	30.90
+Cons@8	28.3	87.3	42.7	62.05
VerifySteer-Uni	35.4	73.8	47.9	9.19
VerifySteer-Bi	37.3	81.0	51.1	9.87
<i>gpt-oss-20B</i>	19.0	90.5	31.4	5.00
<i>Qwen3-14B-thinking</i>	18.5	85.4	30.4	110.87
<i>Llama-3.3-70B</i>	4.4	88.1	8.4	71.35
<i>Qwen2.5-72B</i>	10.1	69.0	17.7	53.03

Table 3: Comparison of VerifySteer with baselines on Hard2Verify. For Qwen3-8B-thinking and Qwen3-14B-thinking, we use the recommended sampling settings (temperature 0.6 and top- p 0.95) and report the mean over 8 runs. Cons@4 and Cons@8 are averaged over 6 self-consistency runs. TFLOPs reports the estimated inference compute per sample.